

Database Maintenance

Updates and Deletes Effect on Table Size

When using PostgreSQL, the size of database tables can grow unexpectedly large with routine UPDATE and DELETE operations.

```
-- Step 1. Generate A New Table
CREATE TABLE rand as (
    SELECT id, random() as score
    FROM generate_series(1, 100000) as id
);

-- Step 2. Check Table Size
SELECT pg_size_pretty(
    pg_total_relation_size('rand')
) as table_size;

/*
+-----+
| table_size |
+-----+
| 4360 kB    |
+-----+
*/

-- Step 3. Update Tuples (~20% of values)
UPDATE rand SET score = 1 where score >
.8;
/* UPDATE 19925 */

-- Step 4. Check Table Size w. ~20% of
tuples updated
SELECT pg_size_pretty(
    pg_total_relation_size('rand')
) as table_size;

/*
+-----+
| table_size |
```

```
+-----+
| 6080 kB |
+-----+
*/
```

PostgreSQL Dead Tuples

In PostgreSQL, when a row is deleted or updated, PostgreSQL creates so-called Dead tuples. Dead tuples are not referenced in the current version of our databases' tables, but still occupy space on disk.

-- Dead tuples contribute to the size of a table but aren't displayed to the DB user: You can check the number of dead tuples with the internal PostgreSQL statistic tables.

```
SELECT
    schemaname,
    relname,
    n_dead_tup
FROM pg_catalog.pg_stat_all_tables
WHERE relname = 'rand';
```

```
/*
+-----+-----+-----+
| schemaname | relname | n_dead_tup |
+-----+-----+-----+
| public     | rand    |      10000 |
+-----+-----+-----+
*/
```

PostgreSQL Vacuuming

In PostgreSQL, to reclaim space from dead tuples, you can use `VACUUM`, `VACUUM ANALYZE`, or `VACUUM FULL`, each comes with a different strategy for clearing dead tuples.

Importance of VACUUM

In PostgreSQL, It's important to occasionally `VACUUM` tables to keep database queries performant and use database space efficiently.

-- Depending on the status of a table's inserts, deletes, and updates, a `\VACUUM`` can reduce space used on disk, or it may

just clear space in the same table for new inserts.

```
-- Before VACUUM
```

```
SELECT
```

```
    schemaname,
```

```
    relname,
```

```
    n_dead_tup
```

```
FROM pg_catalog.pg_stat_all_tables
```

```
/*
```

```
+-----+-----+-----+
```

```
| schemaname | relname | n_dead_tup |
```

```
+-----+-----+-----+
```

```
| public      | rand      |          10000 |
```

```
+-----+-----+-----+
```

```
*/
```

```
VACUUM mocked_data.time_series;
```

```
-- Same Query -> After VACUUM
```

```
/*
```

```
+-----+-----+-----+
```

```
| schemaname | relname | n_dead_tup |
```

```
+-----+-----+-----+
```

```
| public      | rand      |           0 |
```

```
+-----+-----+-----+
```

```
*/
```

PostgreSQL Analyze

In PostgreSQL, `ANALYZE` collects statistics about the contents of tables in the database, and stores the results in the system catalog so PostgreSQL can determine the efficient way to execute a query.

```
-- The statement to analyze a table named
`schema.table`:
ANALYZE schema.table;
```

VACUUM in PostgreSQL

In PostgreSQL, plain `VACUUM` can run in parallel with database operations, but `VACUUM` does not always

```
-- VACUUM `schemaname.tablename` with the
```

fully reduce table sizes. Instead, it marks the space on disk as safe to overwrite with new data.

VACUUM FULL in PostgreSQL

In PostgreSQL, `VACUUM FULL` should be used to fully reclaim database space. However, `VACUUM FULL` rewrites the entire contents of the table into a new location on disk with no extra space allocated. This is an expensive operation and should be used sparingly.

below:



```
VACUUM schemaname.tablename;
```

```
-- Step 1. Check Status of Table - 44.2K
Dead Tuples
```

```
SELECT
    schemaname,
    relname,
    n_dead_tup
FROM pg_catalog.pg_stat_all_tables
WHERE relname = 'rand';

/*
+-----+-----+-----+
| schemaname | relname | n_dead_tup |
+-----+-----+-----+
| public    | rand    |          44157 |
+-----+-----+-----+
*/
```

```
-- Step 2. Run `VACUUM FULL`
VACUUM FULL rand;
```

```
-- Step 3. Confirm dead tuples removed.
```

```
SELECT
    schemaname,
    relname,
    n_dead_tup
FROM pg_catalog.pg_stat_all_tables
WHERE relname = 'rand';

/*
+-----+-----+-----+
| schemaname | relname | n_dead_tup |
+-----+-----+-----+
| public    | rand    |           0 |
+-----+-----+-----+
*/
```

Vacuum and Autovacuum in PostgreSQL

PostgreSQL has a feature called autovacuum, which automatically runs `VACUUM` and `ANALYZE` commands. When enabled, autovacuum checks for tables that have had a large number of inserted, updated or deleted tuples.

-- Autovacuum is enabled on most Database Instances, consider the following:

-- Step 1. Table shows no Autovacuum time

```
select
    schemaname,
    relname,
    last_autovacuum
FROM pg_catalog.pg_stat_all_tables
WHERE relname = 'rand';

/*
+-----+-----+-----+
-+
| schemaname | relname | last_autovacuum
|
+-----+-----+-----+
-+
| public     | rand    |
|
+-----+-----+-----+
-+
*/
```

-- Step 2. Duplicate all rows of a table (`rand`) by re-inserting all rows...

```
INSERT INTO rand (
    SELECT * FROM rand
);
```

-- Step 3. This is a large insert, checking autovacuum time

```
select
    schemaname,
    relname,
    last_autovacuum
FROM pg_catalog.pg_stat_all_tables
WHERE relname = 'rand';
```

```

/*
+-----+-----+-----+
+-----+
| schemaname | relname |
last_autovacuum      |
+-----+-----+-----+
+-----+
| public      | rand      | 2021-07-11
06:24:11.957292-04 |
+-----+-----+-----+
+-----+
*/

```

PostgreSQL Truncate

In PostgreSQL, to improve performance of large deletes, **TRUNCATE** is preferable to **DELETE**, **TRUNCATE** is faster and automatically reclaims the space on disk.

```

-- Step 1. Generate Table
CREATE TABLE rand as (
    SELECT id, random() as score
    FROM generate_series(1, 100000) as id
);

```

```

-- Step 2. Check Table Size
select pg_size_pretty(
    pg_total_relation_size('rand')
) as table_size;

```

```

/*
+-----+
| table_size |
+-----+
| 4360 kB    |
+-----+
*/

```

```

-- Step 3. Option 1: Remove all contents
from table, reclaims no space
DELETE FROM rand WHERE TRUE;

```

```

select pg_size_pretty(
    pg_total_relation_size('rand')
) as table_size;

```

```

) as table_size;

/*

+-----+
| table_size |
+-----+
| 4360 kB    |
+-----+

*/

```

```

-- Step 3. Option 2: With Truncate -
Faster and Automatically Reclaims Space

TRUNCATE rand;

```

```

select pg_size_pretty(
    pg_total_relation_size('rand')
) as table_size;

/*

+-----+
| table_size |
+-----+
| 0 bytes    |
+-----+

*/

```

PostgreSQL All Table Statistics

In PostgreSQL, you can monitor table statistics by querying the view `pg_stat_all_tables`. This view contains statistics like number of dead and live tuples, number of rows inserted, and last vacuum or autovacuum time.

```

-- Sample query to get table statistics
from `pg_stat_all_tables`

SELECT

    schemaname,

    relname,

    n_live_tup,
    n_tup_upd,
    n_tup_del,
    last_vacuum,
    last_autovacuum

FROM pg_catalog.pg_stat_all_tables

WHERE relname = 'clicks';

```

```
/*
+-----+-----+-----+-----+
-----+-----+-----+-----+
-----+-----+
| schemaname | relname | n_live_tup |
n_tup_upd | n_tup_del |
last_vacuum          | last_autovacuum |
+-----+-----+-----+-----+
-----+-----+-----+-----+
-----+-----+
| public      | clicks  |      8200 |
0 |          0 | 2021-06-11
00:23:21.187128-04 |          |
+-----+-----+-----+-----+
-----+-----+-----+-----+
-----+-----+
*/
```



Save



Print



Share ▼